

# **CS:APP Chapter 4**

## **Computer Architecture**

### **Wrap-Up**

**Yuan Tang**

***Adapted from CMU course 15-213***

<http://csapp.cs.cmu.edu>

# Overview

## Wrap-Up of PIPE Design

- Exceptional conditions
- Performance analysis
- Fetch stage design

## Modern High-Performance Processors

- Out-of-order execution

# Exceptions

- Conditions under which processor cannot continue normal operation

## Causes

- Halt instruction (Current)
- Bad address for instruction or data (Previous)
- Invalid instruction (Previous)

## Typical Desired Action

- Complete some instructions
  - Either current or previous (depends on exception type)
- Discard others
- Call exception handler
  - Like an unexpected procedure call



## Our Implementation

- Halt when instruction causes exception

# Exception Examples

## Detect in Fetch Stage

```
jmp $-1                # Invalid jump target

.byte 0xFF             # Invalid instruction code


halt                   # Halt instruction
```

## Detect in Memory Stage

```
irmovl $100,%eax
rmmovl %eax,0x10000(%eax) # invalid address
```

# Exceptions in Pipeline Processor #1

```
# demo-excl1.y  
irmovl $100,%eax  
rmmovl %eax,0x10000(%eax) # Invalid address  
nop  
.byte 0xFF # Invalid instruction code
```



## Desired Behavior

- `rmmovl` should cause exception
- Following instructions should have no effect on processor state

# Exceptions in Pipeline Processor #2

```
# demo-exc2.ys
```

```
0x000:    xorl %eax,%eax    # Set condition codes
```


```
0x002:    jne t            # Not taken
```

```
0x007:    irmovl $1,%eax
```

```
0x00d:    irmovl $2,%edx
```

```
0x013:    halt
```

```
0x014: t: .byte 0xFF      # Target
```

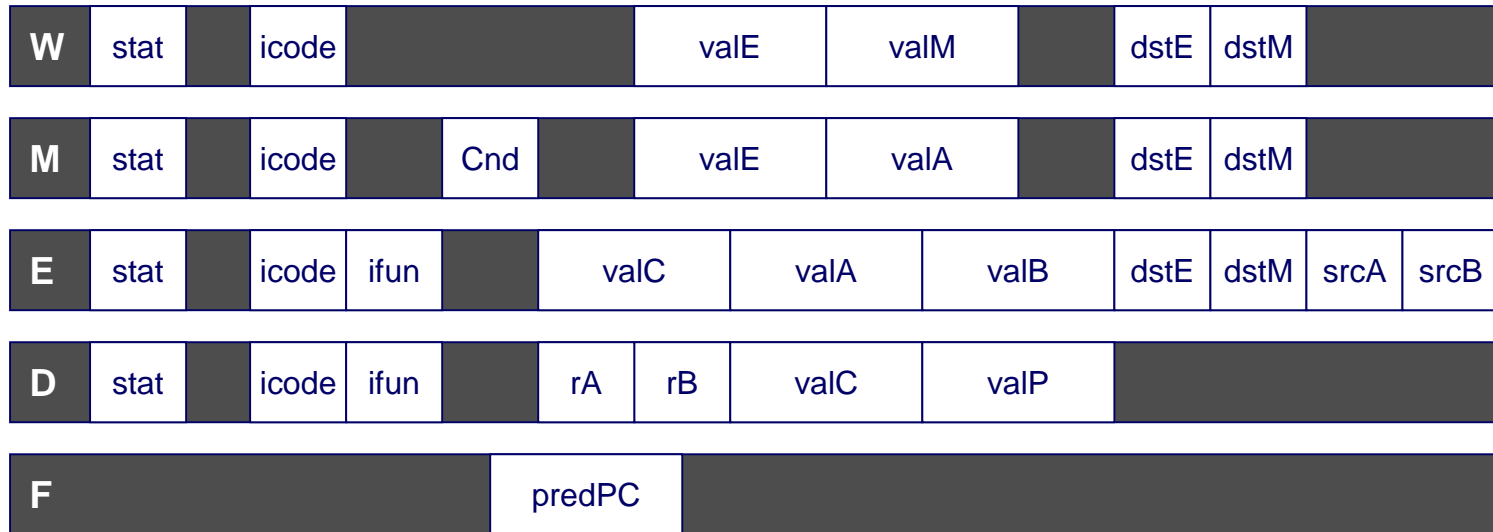


Exception detected

## Desired Behavior

- No exception should occur

# Maintaining Exception Ordering



- Add status field to pipeline registers
- Fetch stage sets to either “AOK,” “ADR” (when bad fetch address), “HLT” (halt instruction) or “INS” (illegal instruction)
- Decode & execute pass values through
- Memory either passes through or sets to “ADR”
- Exception triggered only when instruction hits write back

# Exception Handling Logic

## Fetch Stage

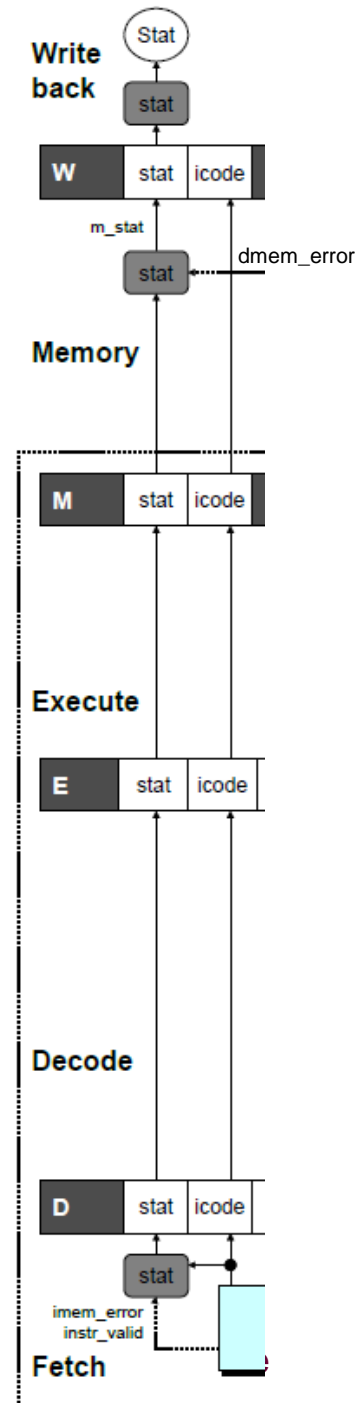
```
# Determine status code for fetched instruction
int f_stat = [
    imem_error: SADR;
    !instr_valid : SINS;
    f_icode == IHALT : SHLT;
    1 : SAOK;
];
```

## Memory Stage

```
# Update the status
int m_stat = [
    dmem_error : SADR;
    1 : M_stat;
];
```

## Writeback Stage

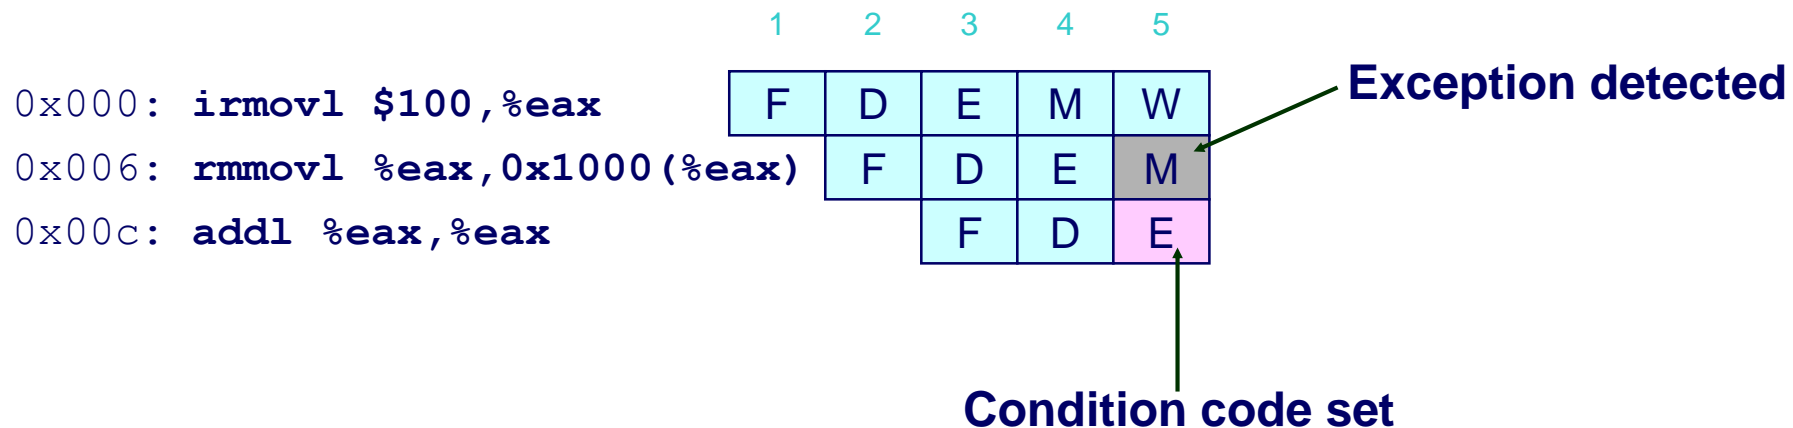
```
int Stat = [
    # SBUB in earlier stages indicates bubble
    W_stat == SBUB : SAOK;
    1 : W_stat;
];
```





# Side Effects in Pipeline Processor

```
# demo-exc3.js
irmovl $100,%eax
rmmovl %eax,0x10000(%eax) # invalid address
addl %eax,%eax           # Sets condition codes
```



## Desired Behavior

- `rmmovl` should cause exception
- No following instruction should have any effect

# Avoiding Side Effects

## Presence of Exception Should Disable State Update

- Invalid instructions are converted to pipeline bubbles
  - Except have stat indicating exception status
- Data memory will not write to invalid address
- Prevent invalid update of condition codes
  - Detect exception in memory stage
  - Disable condition code setting in execute
  - Must happen in same clock cycle
- Handling exception in final stages
  - When detect exception in memory stage
    - » Start injecting bubbles into memory stage on next cycle
  - When detect exception in write-back stage
    - » Stall excepting instruction
- Included in HCL code

# Control Logic for State Changes

## Setting Condition Codes

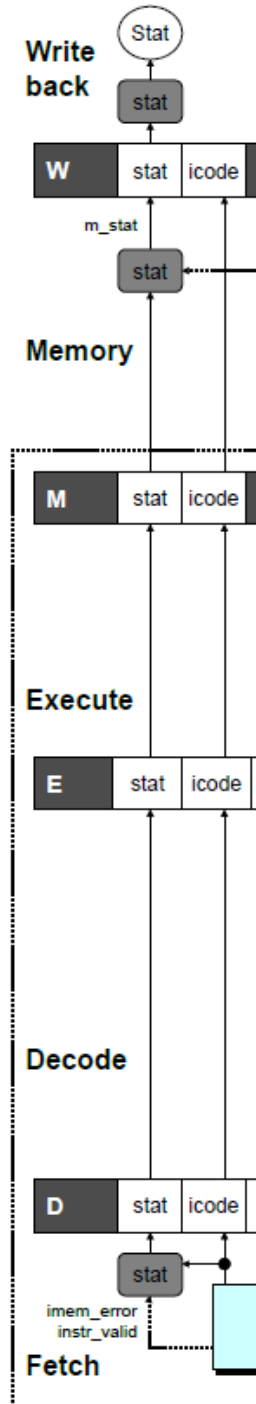
```
# Should the condition codes be updated?
bool set_cc = E_icode == IOPL &&
    # State changes only during normal operation
    !m_stat in { SADR, SINS, SHLT }
    && !W_stat in { SADR, SINS, SHLT };
```

# Stage Control

- Also controls updating of memory

```
# Start injecting bubbles as soon as exception passes
through memory stage
bool M_bubble = m_stat in { SADR, SINS, SHLT }
             || W_stat in { SADR, SINS, SHLT };

# Stall pipeline register W when exception encountered
bool W_stall = W_stat in { SADR, SINS, SHLT };
```



# Rest of Real-Life Exception Handling

## Call Exception Handler

- Push PC onto stack
  - Either PC of faulting instruction or of next instruction
  - Usually pass through pipeline along with exception status
- Jump to handler address
  - Usually fixed address
  - Defined as part of ISA

## Implementation

- Haven't tried it yet!

# Performance Metrics

## Clock rate

- Measured in Gigahertz
- Function of stage partitioning and circuit design
  - Keep amount of work per stage small

## Rate at which instructions executed

- CPI: cycles per instruction
- On average, how many clock cycles does each instruction require?
- Function of pipeline design and benchmark programs
  - E.g., how frequently are branches mispredicted?

# CPI for PIPE

## CPI $\approx$ 1.0

- Fetch instruction each clock cycle
- Effectively process new instruction almost every cycle
  - Although each individual instruction has latency of 5 cycles

## CPI $>$ 1.0

- Sometimes must stall or cancel branches

## Computing CPI

- C clock cycles
- I instructions executed to completion
- B bubbles injected ( $C = I + B$ )

$$\text{CPI} = C/I = (I+B)/I = 1.0 + B/I$$

- Factor  $B/I$  represents average penalty due to bubbles

# CPI for PIPE (Cont.)

$$B/I = LP + MP + RP$$

## Typical Values

### ■ LP: Penalty due to load/use hazard stalling

- Fraction of instructions that are loads 0.25
- Fraction of load instructions requiring stall 0.20
- Number of bubbles injected each time 1

$$\Rightarrow LP = 0.25 * 0.20 * 1 = 0.05$$

### ■ MP: Penalty due to mispredicted branches

- Fraction of instructions that are cond. jumps 0.20
- Fraction of cond. jumps mispredicted 0.40
- Number of bubbles injected each time 2

$$\Rightarrow MP = 0.20 * 0.40 * 2 = 0.16$$

### ■ RP: Penalty due to ret instructions

- Fraction of instructions that are returns 0.02
- Number of bubbles injected each time 3

$$\Rightarrow RP = 0.02 * 3 = 0.06$$

### ■ Net effect of penalties $0.05 + 0.16 + 0.06 = 0.27$

$$\Rightarrow CPI = 1.27 \quad (\text{Not bad!})$$

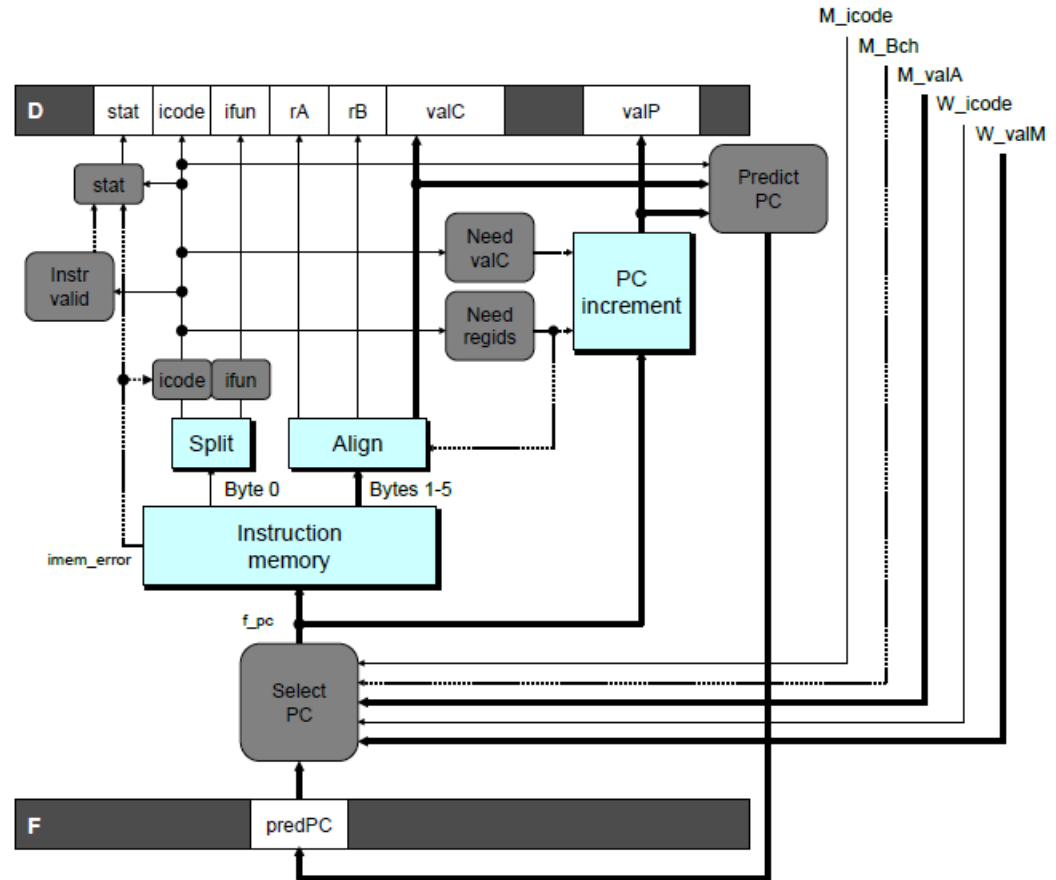
# Fetch Logic Revisited

## During Fetch Cycle

1. Select PC
2. Read bytes from instruction memory
3. Examine icode to determine instruction length
4. Increment PC


## Timing

- Steps 2 & 4 require significant amount of time



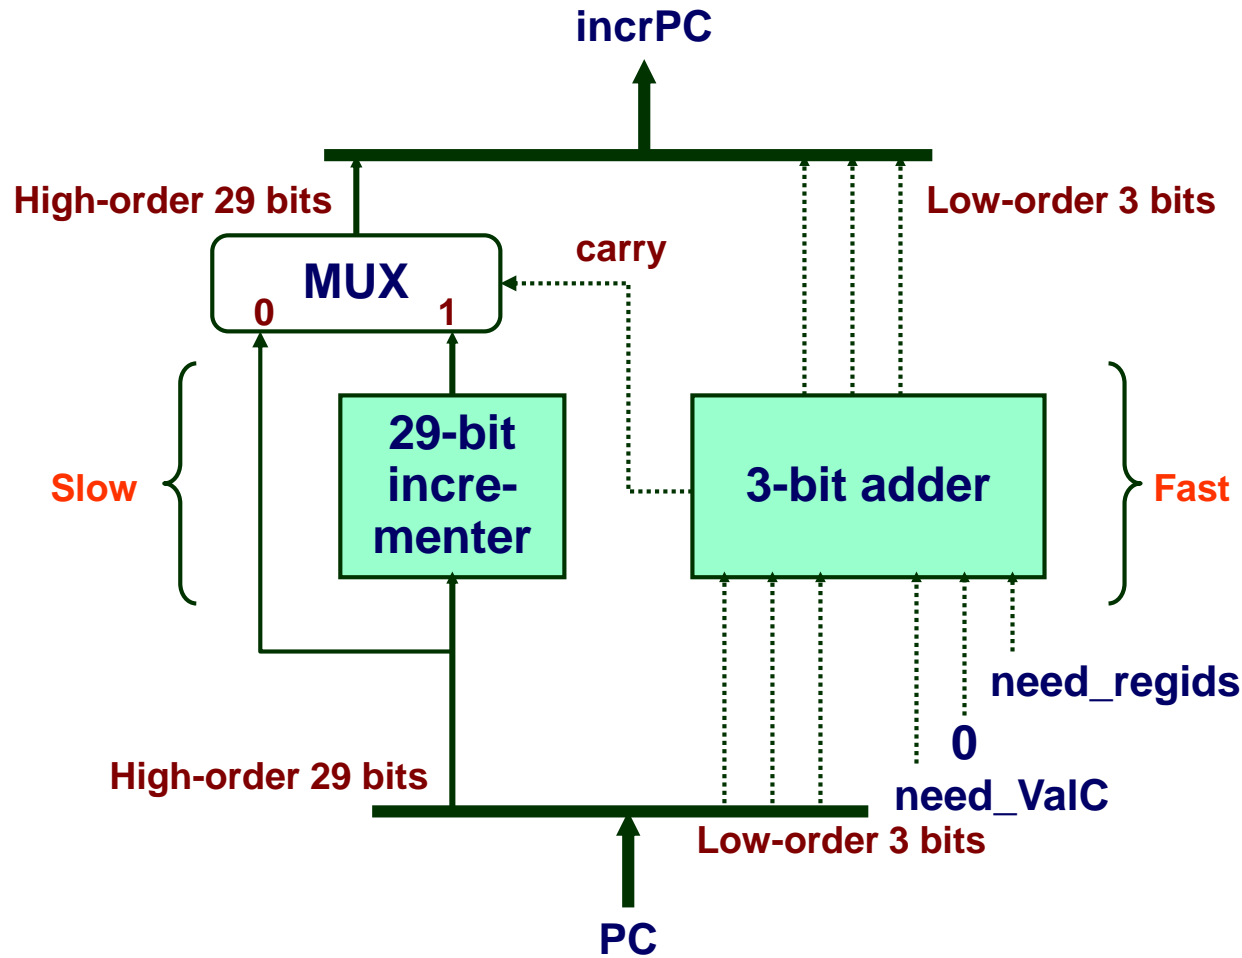


# Standard Fetch Timing

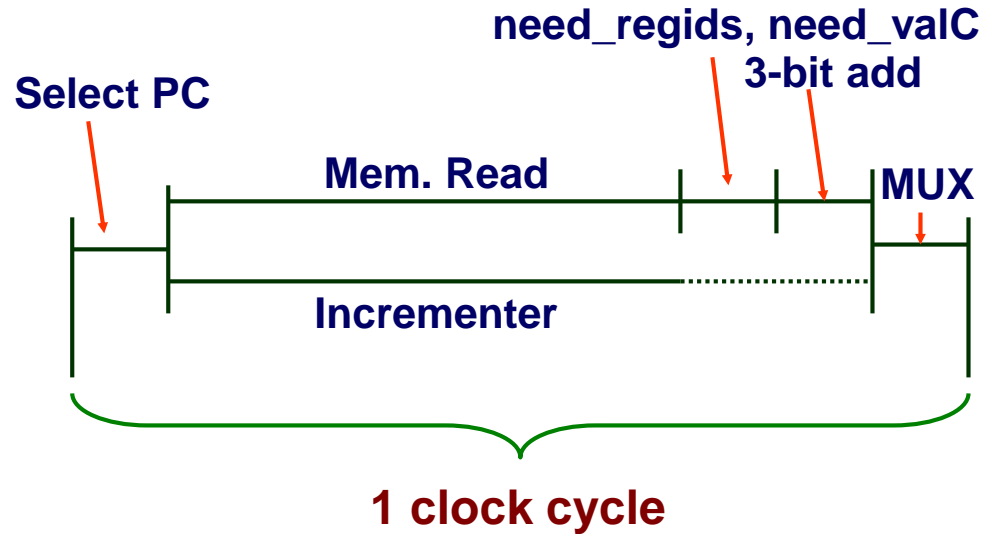


- Must Perform Everything in Sequence
- Can't compute incremented PC until know how much to increment it by

# A Fast PC Increment Circuit



# Modified Fetch Timing

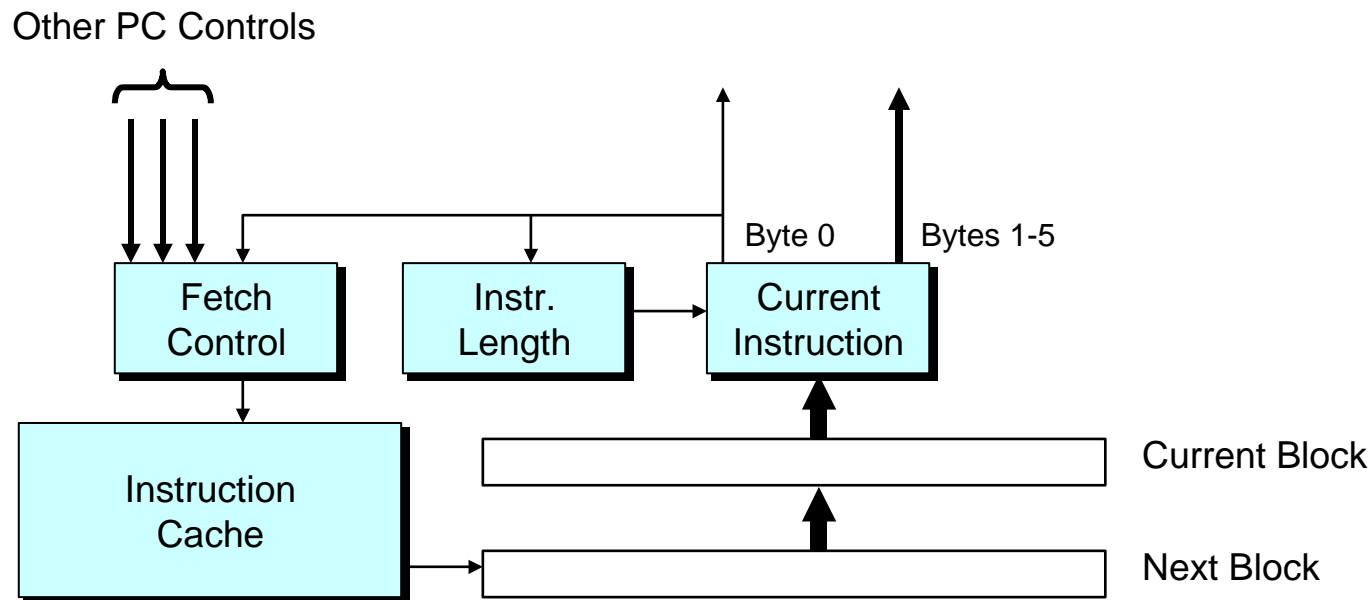


Standard cycle

## 29-Bit Incrementer

- Acts as soon as PC selected
- Output not needed until final MUX
- Works in parallel with memory read

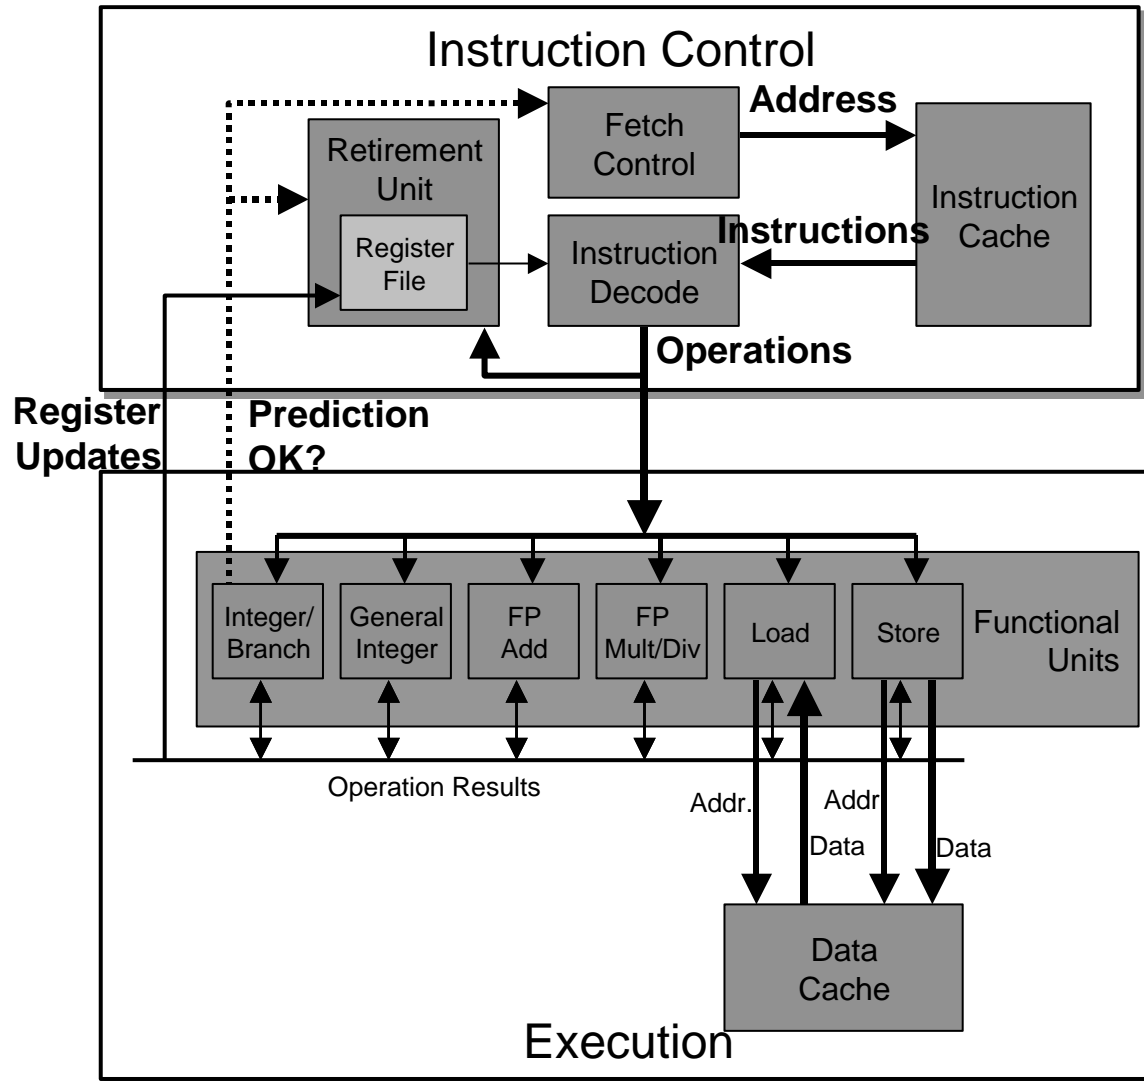
# More Realistic Fetch Logic



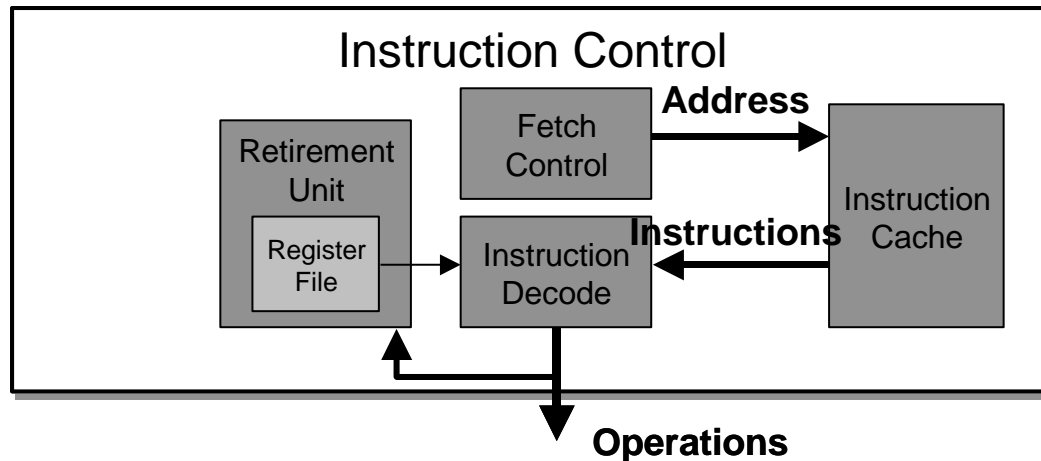
## Fetch Box

- Integrated into instruction cache
- Fetches entire cache block (16 or 32 bytes)
- Selects current instruction from current block
- Works ahead to fetch next block
  - As reaches end of current block
  - At branch target

# Modern CPU Design



# Instruction Control



## Grabs Instruction Bytes From Memory

- Based on Current PC + Predicted Targets for Predicted Branches
- Hardware dynamically guesses whether branches taken/not taken and (possibly) branch target

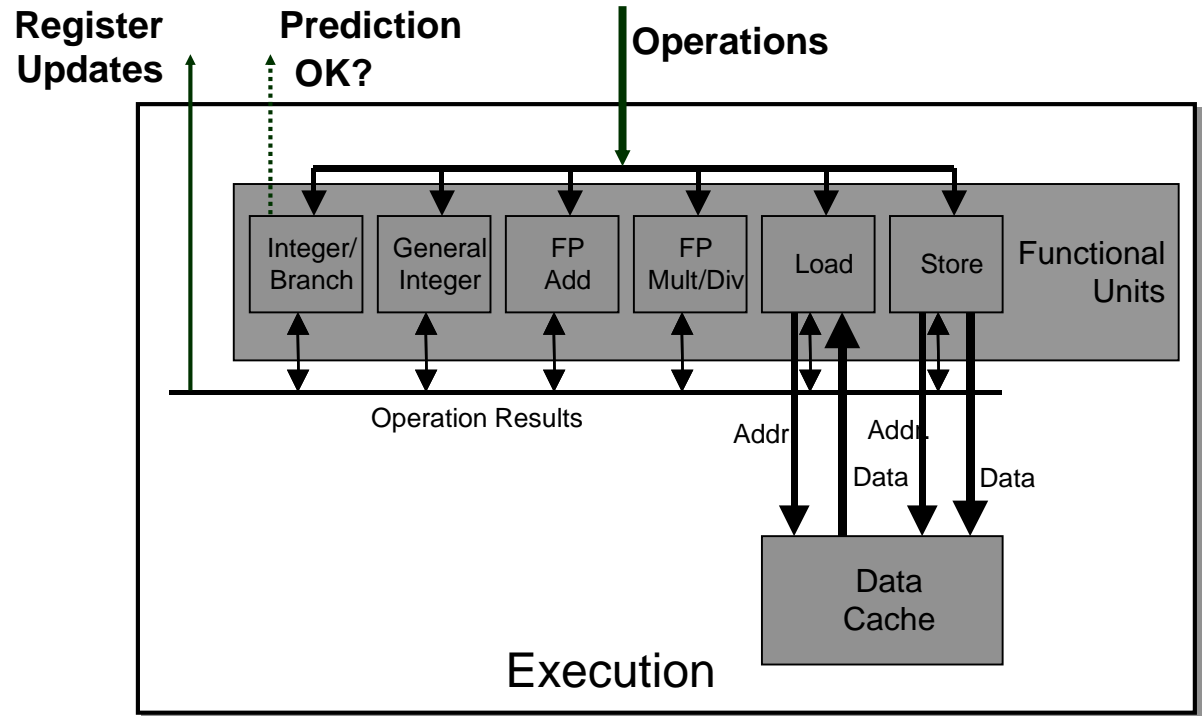
## Translates Instructions Into *Operations*

- Primitive steps required to perform instruction
- Typical instruction requires 1–3 operations

## Converts Register References Into *Tags*

- Abstract identifier linking destination of one operation with sources of later operations

# Execution Unit



- **Multiple functional units**
  - Each can operate independently
- **Operations performed as soon as operands available**
  - Not necessarily in program order
  - Within limits of functional units
- **Control logic**
  - Ensures behavior equivalent to sequential program execution

# CPU Capabilities of Intel iCore7

## Multiple Instructions Can Execute in Parallel

- 1 load
- 1 store
- 1 FP multiplication or division
- 1 FP addition
- > 1 integer operation

## Some Instructions Take > 1 Cycle, but Can be Pipelined

■ Instruction	Latency	Cycles/Issue
■ Load / Store	3	1
■ Integer Multiply	3	1
■ Integer Divide	11—21	5—13
■ Double/Single FP Multiply	4	1
■ Double/Single FP Add	3	1
■ Double/Single FP Divide	10—15	6—11



# iCore Operation

**Translates instructions dynamically into “Uops”**

- ~118 bits wide
- Holds operation, two sources, and destination

**Executes Uops with “Out of Order” engine**

- Uop executed when
  - Operands available
  - Functional unit available
- Execution controlled by “Reservation Stations”
  - Keeps track of data dependencies between uops
  - Allocates resources

# High-Performance Branch Prediction

## Critical to Performance

- Typically 11–15 cycle penalty for misprediction

## Branch Target Buffer

- 512 entries
- 4 bits of history
- Adaptive algorithm
  - Can recognize repeated patterns, e.g., alternating taken–not taken


## Handling BTB misses

- Detect in ~cycle 6
- Predict taken for negative offset, not taken for positive
  - Loops vs. conditionals

# Example Branch Prediction

## Branch History

- Encode information about prior history of branch instructions
- Predict whether or not branch will be taken



## State Machine

- Each time branch taken, transition to right
- When not taken, transition to left
- Predict branch taken when in state Yes! or Yes?

# Processor Summary

## Design Technique

- Create uniform framework for all instructions
  - Want to share hardware among instructions
- Connect standard logic blocks with bits of control logic

## Operation

- State held in memories and clocked registers
- Computation done by combinational logic
- Clocking of registers/memories sufficient to control overall behavior

## Enhancing Performance

- Pipelining increases throughput and improves resource utilization
- Must make sure to maintain ISA behavior